

Determining majority in networks with local interactions and very small local memory

George B. Mertzios¹ Sotiris E. Nikolettseas²
Christoforos L. Raptopoulos^{2,3} Paul G. Spirakis^{2,4}

¹School of Engineering and Computing Sciences, Durham University, UK

²Computer Technology Institute (CTI) and University of Patras, Greece

³Computer Science Department, University of Geneva, Switzerland

⁴Department of Computer Science, University of Liverpool, UK

ICALP 2014

Algorithms and Complexity in Durham (ACiD) Seminar,
Durham University, November 2014

Consensus in distributed systems

In distributed systems:

- a collection of n independent **entities** (or **nodes**)
- entities **interact** / exchange **messages** to coordinate their actions
- interactions must satisfy some constraints, e.g.:
 - synchronous vs. asynchronous,
 - not every entity can interact with all others (network structure),
 - how often two specific entities may interact, etc.

Consensus in distributed systems

In distributed systems:

- a collection of n independent **entities** (or **nodes**)
- entities **interact** / exchange **messages** to coordinate their actions
- interactions must satisfy some constraints, e.g.:
 - synchronous vs. asynchronous,
 - not every entity can interact with all others (network structure),
 - how often two specific entities may interact, etc.

A central problem in distributed systems:

Definition (Consensus)

Let each node have an input value. A solution for the **consensus** problem must guarantee:

- **Termination:** every node eventually decides on some value,
- **Agreement:** all nodes decide on the same value,
- **Validity:** the decided value must be the input of some node.

Consensus in distributed systems

Many applications of the consensus problem, e.g.:

- leader election
- distributed ranking [Jung et al., *ISIT*, 2012]

The **majority** problem:

- a natural special case of the consensus problem
- the **agreed value** must be the input value of the **majority** of the nodes
- two or more different input values (or colors)
[Angluin et al., *Distributed Computing*, 2008]
[Becchetti et al., *SPAA*, 2014]
- many applications, e.g.:
 - voting [Kearns et al., *WINE*, 2008]
 - epidemiology and interacting particle systems
[Liggett, *Interacting Particle Systems*, 2004]
 - social networks [Mizrachi, *MSc thesis*, Ben-Gurion University, 2013]
[Mossel et al., *Auton. Agents & Multi-Agent Systems*, 2014]

Computing the majority

- To solve the majority problem in a network:
 - we need assumptions on the model of computation
 - In the “traditional” settings: “strong” models
 - **central** authority, **unlimited** memory, **full** information about the network
 - efficiently computable
 - the goal is to minimize the number of comparisons
- [Saks et al., *Combinatorica*, 1991]
[De Marco et al., *Combinatorics, Probability and Computing*, 2006]

Computing the majority

- To solve the majority problem in a network:
 - we need assumptions on the model of computation
- In the “traditional” settings: “strong” models
 - **central** authority, **unlimited** memory, **full** information about the network
 - efficiently computable
 - the goal is to minimize the number of comparisons
[Saks et al., *Combinatorica*, 1991]
[De Marco et al., *Combinatorics, Probability and Computing*, 2006]
- In “modern” settings: “weaker” models
 - **no** central authority, **limited** memory, **partial** or **no** information
 - a node does not know:
 - its own identity
 - the identities of the nodes it can interact with (i.e. its neighbors)
 - when it will interact with other nodes
 - one way to model such systems is using **population protocols**

Population protocols

- Population V of $|V| = n$ entities (i.e. nodes)
- A **population protocol** \mathcal{A} consists of:
 - finite **input** and **output alphabets** X and Y
 - a finite set of **states** Q
 - an **input function** $I : X \rightarrow Q$
 - an **output function** $O : Q \rightarrow Y$
 - a **transition function** $\delta : Q \times Q \rightarrow Q \times Q$

Population protocols

- Population V of $|V| = n$ entities (i.e. nodes)
- A **population protocol** \mathcal{A} consists of:
 - finite **input** and **output alphabets** X and Y
 - a finite set of **states** Q
 - an **input function** $I : X \rightarrow Q$
 - an **output function** $O : Q \rightarrow Y$
 - a **transition function** $\delta : Q \times Q \rightarrow Q \times Q$
- The **result** of an interaction between nodes u and v depends **only** on their **states** q_u and q_v

Population protocols

- Population V of $|V| = n$ entities (i.e. nodes)
- A **population protocol** \mathcal{A} consists of:
 - finite **input** and **output alphabets** X and Y
 - a finite set of **states** Q
 - an **input function** $I : X \rightarrow Q$
 - an **output function** $O : Q \rightarrow Y$
 - a **transition function** $\delta : Q \times Q \rightarrow Q \times Q$
- The **result** of an interaction between nodes u and v depends **only** on their **states** q_u and q_v
- A population protocol is **symmetric** if interactions have **no “direction”**:
 - $\delta(q_u, q_v) = (q'_u, q'_v) \iff \delta(q_v, q_u) = (q'_v, q'_u),$
for every pair of states $q_u, q_v \in Q$

Population protocols

- Population V of $|V| = n$ entities (i.e. nodes)
- A **population protocol** \mathcal{A} consists of:
 - finite **input** and **output alphabets** X and Y
 - a finite set of **states** Q
 - an **input function** $I : X \rightarrow Q$
 - an **output function** $O : Q \rightarrow Y$
 - a **transition function** $\delta : Q \times Q \rightarrow Q \times Q$
- The **result** of an interaction between nodes u and v depends **only** on their **states** q_u and q_v
- A population protocol is **symmetric** if interactions have **no “direction”**:
 - $\delta(q_u, q_v) = (q'_u, q'_v) \iff \delta(q_v, q_u) = (q'_v, q'_u),$
for every pair of states $q_u, q_v \in Q$
- Otherwise, for every interaction, one of the nodes is the **initiator**

Population protocols

Schedulers

Terminology:

- The interaction **order** is chosen by an **adversary** (scheduler)
- To allow meaningful computations: **scheduler** must be **fair**
 - we do not allow avoidance of a possible step forever
 - for any two state configurations C_1, C_2 , where C_2 is reachable from C_1 :
if C_1 occurs **infinitely** often $\Rightarrow C_2$ also occurs **infinitely** often

Terminology:

- The interaction **order** is chosen by an **adversary** (scheduler)
- To allow meaningful computations: **scheduler** must be **fair**
 - we do not allow avoidance of a possible step forever
 - for any two state configurations C_1, C_2 , where C_2 is reachable from C_1 : if C_1 occurs **infinitely** often $\Rightarrow C_2$ also occurs **infinitely** often
- The **interaction graph** $G = (V, E)$ of the population:
 - the entities of the population are arranged on the nodes V
 - only neighboring nodes are allowed to interact
- The **probabilistic scheduler**:
 - a special case of a fair scheduler
 - **directed** case: every directed edge (u, v) is chosen uniformly at random (u is the initiator)
 - **undirected** case: replace edge $\{u, v\}$ by directed edges $(u, v), (v, u)$

Terminology:

Definition

Given the **probabilistic scheduler**, a population protocol \mathcal{A} **computes** a **function** g with **error probability** ε if for every input configuration C_0 the population eventually reaches a configuration C such that with **probability** at least $1 - \varepsilon$:

- (a) **all nodes** have **output** $g(C_0)$
- (b) this remains true for **any configuration** reachable from C

Population protocols

Computation

Terminology:

Definition

Given the **probabilistic scheduler**, a population protocol \mathcal{A} **computes** a **function** g with **error probability** ε if for every input configuration C_0 the population eventually reaches a configuration C such that with **probability** at least $1 - \varepsilon$:

- (a) **all nodes** have **output** $g(C_0)$
- (b) this remains true for **any configuration** reachable from C

Definition

A population protocol \mathcal{A} **stably computes** a **function** g if for **every fair scheduler** the population eventually reaches a configuration C that satisfies both (a) and (b).

Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**


Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**

Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**


Example:



Population protocols for computing the majority

- Computing the majority in distributed settings has been mainly studied in **homogeneous** populations (i.e. the **complete graph**)
- The following simple 3-state population protocol was introduced in [Angluin et al., *Distributed Computing*, 2008]
 - initially nodes have 2 possible states: **r** and **g**
 - during the execution, a node can have 3 possible states: **r**, **g**, and **b**
 - interactions are dictated by the **probabilistic scheduler**
 - the **3×3 transition table** can be summarized as follows:
 - node u of state **r** “hits” node v of state **g** $\Rightarrow v$ comes to state **b**
 - node u of state **g** “hits” node v of state **r** $\Rightarrow v$ comes to state **b**
 - node u of state **r** / **g** “hits” node v of state **b** $\Rightarrow v$ comes to state **r** / **g**

Example:



Population protocols for computing the majority

- In the protocol of [Angluin et al., *Distributed Computing*, 2008]:
 - if the underlying **interaction graph** is **complete** (with n vertices)
 - and the **initial difference** between majority and minority is $\Omega(\sqrt{n \log n})$
 - then it **converges** to the initial majority in $O(n \log n)$ time w.h.p.
- A similar protocol for the complete graph has been studied in [Perron et al., *INFOCOM*, 2009]

Population protocols for computing the majority

- In the protocol of [Angluin et al., *Distributed Computing*, 2008]:
 - if the underlying interaction graph is complete (with n vertices)
 - and the initial difference between majority and minority is $\Omega(\sqrt{n \log n})$
 - then it converges to the initial majority in $O(n \log n)$ time w.h.p.
- A similar protocol for the complete graph has been studied in [Perron et al., *INFOCOM*, 2009]

In the case of arbitrary interaction graphs:

- how fast can such protocols terminate?
- do they compute the correct initial majority with high probability?
- is it possible to compute majority with probability 1?
- how many states (per node) do we need to compute majority?
- how large should be the difference between initial majority / minority?

Our results

First result: the ambassador protocol

Theorem

- There exists a *4-state* protocol, the *ambassador protocol*, which *stably computes* the initial majority value:
 - for *any* interaction graph G ,
 - for *any* initial difference between majority / minority,
 - with *probability 1*.
- There does *not* exist any *3-state* protocol with these properties

Our results

First result: the ambassador protocol

Theorem

- There exists a *4-state* protocol, the *ambassador protocol*, which *stably computes* the initial majority value:
 - for *any* interaction graph G ,
 - for *any* initial difference between majority / minority,
 - with *probability 1*.
- There does *not* exist any *3-state* protocol with these properties

Theorem

Under the *probabilistic scheduler*:

- The *4-state ambassador* protocol runs in *expected polynomial time*.
- If the interaction graph G is *complete* and the *initial difference* is $\Theta(n)$, then the protocol terminates in expected time $O(n \log n)$.

Our results

Second result: a detailed analysis of the protocol of *Angluin et al.* on an arbitrary interaction graph G (under the **probabilistic scheduler**)

Theorem

If the types **r** and **g** are distributed **uniformly at random** on the vertices of G , the protocol **converges** to the **initial majority** with **probability** $\geq \frac{1}{2}$.

Our results

Second result: a detailed analysis of the protocol of *Angluin et al.* on an arbitrary interaction graph G (under the **probabilistic scheduler**)

Theorem

If the types **r** and **g** are distributed **uniformly at random** on the vertices of G , the protocol **converges** to the **initial majority** with **probability** $\geq \frac{1}{2}$.

Theorem

There exists an **infinite family** $\{G_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol **fails** with **high probability**, even when the **initial difference** between majority / minority is $n - \Theta(\log n)$.

Our results

Second result: a detailed analysis of the protocol of *Angluin et al.* on an arbitrary interaction graph G (under the **probabilistic scheduler**)

Theorem

If the types **r** and **g** are distributed **uniformly at random** on the vertices of G , the protocol **converges** to the **initial majority** with **probability** $\geq \frac{1}{2}$.

Theorem

There exists an **infinite family** $\{G_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol **fails** with **high probability**, even when the **initial difference** between majority / minority is $n - \Theta(\log n)$.

Theorem

There exists an **infinite family** $\{G'_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol terminates in **exponential expected time**.

The 4-state ambassador protocol

The **symmetric** 4-state **ambassador** protocol:

- every node always has a color (**r** or **g**)
 - every node may (or may not) have an extra token (called **ambassador**)
- ⇒ every node has 4 possible states: (**r**,0), (**r**,1), (**g**,0), (**g**,1)
- having an ambassador, a node can promote its color to a neighbor
 - initially every node is at state (**r**,1) or (**g**,1), i.e. with an ambassador

The 4-state ambassador protocol

The **symmetric** 4-state **ambassador protocol**:

- every node always has a color (**r** or **g**)
 - every node may (or may not) have an extra token (called **ambassador**)
- ⇒ every node has 4 possible states: (**r**,0), (**r**,1), (**g**,0), (**g**,1)
- having an ambassador, a node can promote its color to a neighbor
 - initially every node is at state (**r**,1) or (**g**,1), i.e. with an ambassador

When two nodes u and v interact, then:

- if **both** u and v have an **ambassador**:
 - if u and v have the **same** color, **nothing** happens
 - if u and v have **different** color, they both **lose** their **ambassadors**

The 4-state ambassador protocol

The **symmetric** 4-state **ambassador protocol**:

- every node always has a color (**r** or **g**)
 - every node may (or may not) have an extra token (called **ambassador**)
- ⇒ every node has 4 possible states: (**r**,0), (**r**,1), (**g**,0), (**g**,1)
- having an ambassador, a node can promote its color to a neighbor
 - initially every node is at state (**r**,1) or (**g**,1), i.e. with an ambassador

When two nodes u and v interact, then:

- if **both** u and v have an **ambassador**:
 - if u and v have the **same** color, **nothing** happens
 - if u and v have **different** color, they both **lose** their **ambassadors**
- if u has an **ambassador** and v does not:
 - the **ambassador** of u **moves** to v
 - v takes the color of u

The 4-state ambassador protocol

The **symmetric** 4-state **ambassador protocol**:


- every node always has a color (**r** or **g**)
 - every node may (or may not) have an extra token (called **ambassador**)
- ⇒ every node has 4 possible states: (**r**,0), (**r**,1), (**g**,0), (**g**,1)
- having an ambassador, a node can promote its color to a neighbor
 - initially every node is at state (**r**,1) or (**g**,1), i.e. with an ambassador

When two nodes u and v interact, then:

- if **both** u and v have an **ambassador**:
 - if u and v have the **same** color, **nothing** happens
 - if u and v have **different** color, they both **lose** their **ambassadors**
- if u has an **ambassador** and v does not:
 - the **ambassador** of u **moves** to v
 - v takes the color of u
- if **neither** u **nor** v have an ambassador:
 - **nothing** happens


The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:




The 4-state ambassador protocol

Example:



The 4-state ambassador protocol

Example:




For **any fair scheduler**:

- the ambassadors of the minority will eventually all die out
- the remaining ambassadors will eventually color all the graph

The 4-state ambassador protocol

Example:



For **any fair scheduler**:

- the ambassadors of the minority will eventually all die out
- the remaining ambassadors will eventually color all the graph

Theorem (correctness)

- The 4-state ambassador protocol **stably computes** the initial majority:
 - for **any** interaction graph G ,
 - for **any** initial difference between majority / minority,
 - with **probability 1**.

Lower bound on the number of states

Theorem

Let P be a population protocol that *stably computes* the *majority* function in an *arbitrary* 2-type population and for an *arbitrary* interaction graph. Then P needs *at least 4 states*.

Lower bound on the number of states

Theorem

Let P be a population protocol that *stably computes* the *majority* function in an *arbitrary* 2-type population and for an *arbitrary* interaction graph. Then P needs *at least 4 states*.

Proof (sketch, by contradiction).

- Assume P has 3 states \mathbf{r} , \mathbf{g} , \mathbf{b}
- For at least one of the two input colors (say \mathbf{r}):
 - starting with a majority of \mathbf{r} ,
 - eventually all nodes have the same state $q \in \{\mathbf{r}, \mathbf{g}, \mathbf{b}\}$
- We construct two instances C_1, C_2 on the same population such that:
 - C_1 and C_2 have different initial majorities
 - there exists a fair scheduler that brings both C_1 and C_2 to the same intermediate configuration
 - contradiction



The 4-state ambassador protocol

For the **probabilistic scheduler**:

Theorem

If $\Delta > 0$ is the *initial difference* between majority / minority, the 4-state ambassador protocol converges in expected:

- $O(n^6)$ time for an *arbitrary* connected graph G
- $O\left(\frac{\ln n}{\Delta} n^2\right)$ time for the *complete graph* K_n .

Proof based on:

- random walks on graphs and coupon collector arguments

Therefore:

- in the **complete graph** K_n , when $\Delta = \omega(\sqrt{n} \log n)$, the ambassador protocol converges in expected $O(n\sqrt{n})$ time
- **a bit slower** than $O(n \log n)$ of the 3-state protocol of [Angluin et al., *Distributed Computing*, 2008]
- but **always correct**

The protocol of Angluin et al. in arbitrary graphs

Assuming the probabilistic scheduler:

- What can we achieve with a 3-state protocol?
 - it cannot stably compute majority on arbitrary graphs
 - but it might compute majority with large enough probability.

The 3-state protocol of Angluin et al.:

- Converges quickly to the correct initial majority whp in the clique (for sufficiently large majority).
- What about arbitrary graphs?

The protocol of Angluin et al. in arbitrary graphs

Assuming the **probabilistic scheduler**:

- What can we achieve with a **3-state** protocol?
 - it cannot stably compute majority on arbitrary graphs
 - but it **might** compute majority with large enough probability.

The **3-state protocol of Angluin et al.**:

- Converges **quickly** to the **correct** initial majority whp in the clique (for sufficiently large majority).
- What about **arbitrary** graphs?

Theorem

*If the types **r** and **g** are distributed **uniformly at random** on the vertices of G , the protocol **converges** to the **initial majority** with **probability** $\geq \frac{1}{2}$.*

- Proof based on Hall's Marriage Theorem.

The protocol of Angluin et al. in arbitrary graphs

- The model of Angluin et al. can be abstracted by a **Markov chain** \mathcal{M} :
 - \mathcal{M} has **states** (R_t, G_t) , where R_t (resp. G_t) is the set of nodes of type **r** (resp. **g**) at time t
 - **symmetries** of the interaction graph can reduce the size of the state space; e.g. in the clique K_n , the set of states is just $(|R_t|, |G_t|)$.
 - The analysis of \mathcal{M} on arbitrary graphs is complicated; for the clique exact formulae can be found [Perron et al., *INFOCOM*, 2009].

The protocol of Angluin et al. in arbitrary graphs

- The model of Angluin et al. can be abstracted by a **Markov chain** \mathcal{M} :
 - \mathcal{M} has **states** (R_t, G_t) , where R_t (resp. G_t) is the set of nodes of type **r** (resp. **g**) at time t
 - **symmetries** of the interaction graph can reduce the size of the state space; e.g. in the clique K_n , the set of states is just $(|R_t|, |G_t|)$.
 - The analysis of \mathcal{M} on arbitrary graphs is complicated; for the clique exact formulae can be found [Perron et al., *INFOCOM*, 2009].
- We define 2 stochastic processes that **filter** the information from \mathcal{M} :

Definition (The **Blank Process** \mathcal{W})

$$\mathcal{W}(t) \stackrel{\text{def}}{=} \langle \# \text{ nodes of type } \mathbf{b} \text{ at time } t \rangle$$

The protocol of Angluin et al. in arbitrary graphs

Definition (The Contest Process \mathcal{C})

- We recursively **pair** the **state changing** transitions in \mathcal{M} as follows:
 - each transition that **increases** the **blanks** ($\mathbf{g} \rightarrow \mathbf{r}$ or $\mathbf{r} \rightarrow \mathbf{g}$)
 - with the **earliest subsequent** transition that **decreases** the **blanks** ($\mathbf{g} \rightarrow \mathbf{b}$ or $\mathbf{r} \rightarrow \mathbf{b}$) and is **not paired yet**.
- define $\tau(t) \stackrel{\text{def}}{=} \langle \# \text{ pairs until time } t \rangle$
- \mathcal{C} is defined over time scale τ
- Initially set $\mathcal{C}(0) = |R_0|$, and recursively:

$$\mathcal{C}(\tau) = \begin{cases} \mathcal{C}(\tau - 1) + 1, & \text{if } \tau\text{-th pair is } (\mathbf{r} \rightarrow \mathbf{g}, \mathbf{r} \rightarrow \mathbf{b}) \\ \mathcal{C}(\tau - 1) - 1, & \text{if } \tau\text{-th pair is } (\mathbf{g} \rightarrow \mathbf{r}, \mathbf{g} \rightarrow \mathbf{b}) \text{ and} \\ \mathcal{C}(\tau - 1), & \text{otherwise.} \end{cases}$$


The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r
2	0	1	2	b \rightarrow r

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r
2	0	1	2	b \rightarrow r

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:




t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r
2	0	1	2	b \rightarrow r
3	1	0	1	g \rightarrow b

The protocol of Angluin et al. in arbitrary graphs

- The Contest Process keeps track of the **battle** between **g** and **r**
- $\mathcal{C}(\tau)$ counts the number of:
 - nodes of type **r** and
 - nodes of type **b** that were previously of type **r**

Example:



t	$\tau = \tau(t)$	$\mathcal{W}(t)$	$\mathcal{C}(\tau)$	transitions
0	0	0	2	-
1	0	1	2	g \rightarrow r
2	0	1	2	b \rightarrow r
3	1	0	1	g \rightarrow b

The protocol of Angluin et al. in arbitrary graphs

- \mathcal{W} and \mathcal{C} are **dependent** and **not Markov chains**
- \mathcal{C} is defined on **different time scale** than \mathcal{W} and \mathcal{M}
- \mathcal{W} decreases \Rightarrow pair of transitions in $\mathcal{M} \Rightarrow$ transition step in \mathcal{C}
- Under assumptions on $|R_t|$ and $|G_t|$, we can **dominate** both \mathcal{W} and \mathcal{C} in the clique by appropriate **birth-death processes**
- Combining the above, we can prove that under the probabilistic scheduler the protocol of Angluin et al. in the **clique** is **robust**:

Theorem

For every constant $\epsilon < 1/7$ in the **complete graph** K_n :

- if we initially have **at most ϵn** type **r** nodes
- then the probability that the **minority r** wins is **exponentially small** in n .

The protocol of Angluin et al. in arbitrary graphs

Convergence to minority whp

Theorem

There exists an *infinite family* $\{G_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol *fails* with *high probability*, even when the *initial difference* between majority / minority is $n - \Theta(\log n)$.

The protocol of Angluin et al. in arbitrary graphs


Convergence to minority whp

Theorem

There exists an *infinite family* $\{G_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol *fails* with *high probability*, even when the *initial difference* between majority / minority is $n - \Theta(\log n)$.

Proof (sketch).

- Let $n_1 \geq 100 \ln n$ and consider the *lollipop graph*:
 - line L_{n-n_1} with leftmost vertex u connected to vertex v of clique K_{n_1}
 - $L_{n-n_1} \cup \{v\}$ is of type **r** and $K_{n_1} \setminus \{v\}$ is of type **g**



The protocol of Angluin et al. in arbitrary graphs

Convergence to minority whp (cntd.)

Proof sketch. (cntd.)

- Define similarly Blank and Contest processes \mathcal{W}' and \mathcal{C}' on K_{n_1}
- These are **slightly different** than before, because of the edge $\{u, v\}$.
- Using \mathcal{W}' and \mathcal{C}' we first show that:

$$\Pr(\text{all } K_{n_1} \text{ becomes } \mathbf{r}) = e^{-\Omega(n_1)}$$

The protocol of Angluin et al. in arbitrary graphs

Convergence to minority whp (cntd.)

Proof sketch. (cntd.)

- Define similarly Blank and Contest processes \mathcal{W}' and \mathcal{C}' on K_{n_1}
- These are **slightly different** than before, because of the edge $\{u, v\}$.
- Using \mathcal{W}' and \mathcal{C}' we first show that:

$$\Pr(\text{all } K_{n_1} \text{ becomes } \mathbf{r}) = e^{-\Omega(n_1)}$$

- Second, we prove that in a line L_{n-n_1} with a **single** vertex of type **g** and the rest of type **r**:

$$\Pr(\text{all } L_{n-n_1} \text{ becomes } \mathbf{g}) = \Omega\left(\frac{1}{n-n_1}\right)$$

The protocol of Angluin et al. in arbitrary graphs

Convergence to minority whp (cntd.)

Proof sketch. (cntd.)

- Define similarly Blank and Contest processes \mathcal{W}' and \mathcal{C}' on K_{n_1}
- These are **slightly different** than before, because of the edge $\{u, v\}$.
- Using \mathcal{W}' and \mathcal{C}' we first show that:

$$\Pr(\text{all } K_{n_1} \text{ becomes } \mathbf{r}) = e^{-\Omega(n_1)}$$

- Second, we prove that in a line L_{n-n_1} with a **single** vertex of type **g** and the rest of type **r**:

$$\Pr(\text{all } L_{n-n_1} \text{ becomes } \mathbf{g}) = \Omega\left(\frac{1}{n-n_1}\right)$$

- The above imply that, for $n_1 \geq 100 \ln n$, the minority **g** in the clique K_{n_1} has **enough attempts** to take over the whole graph.




The protocol of Angluin et al. in arbitrary graphs

Exponential expected convergence time

Theorem

There exists an *infinite family* $\{G'_n\}_{n \in \mathbb{N}}$ of interaction graphs where the protocol terminates in *exponential expected time*.


- We consider the family of graphs consisting of a clique K_{n_1} of type **g** and a clique K_{n_2} of type **r**, connected with an edge.



- The proof builds upon the proof ideas for the robustness of the protocol in the clique.

The protocol of Angluin et al. in arbitrary graphs

Exponential expected convergence time




Main idea:

- if **vertex** v becomes **r**:
 - K_{n_1} needs expected exponential time in n_1 to become of type **r**

The protocol of Angluin et al. in arbitrary graphs

Exponential expected convergence time



Main idea:

- if **vertex** v becomes **r**:
 - K_{n_1} needs expected exponential time in n_1 to become of type **r**
- if **vertex** u becomes **g**:
 - K_{n_2} needs expected exponential time in n_2 to become of type **g**

Summary and Open Problems

- A 4-state symmetric (ambassador) protocol that always computes the majority
 - this is not possible with 3 states per node
- A detailed analysis of the majority protocol of Angluin et al. on arbitrary graphs
 - although it converges correctly and quickly whp in the clique,
 - this is not the case for arbitrary graphs

Summary and Open Problems

Open problems:

- A “good” 3-state protocol for majority on arbitrary graphs (under the probabilistic scheduler) ?
- Other computations than majority ?
 - average value, median, ...
- What can we compute by allowing more powerful agents?
 - Non-deterministic interactions
 - More memory; what kind of functions can we (stably) compute with (say) 10, or $\log \log n$ states per vertex?
- What if every interaction involves more than 2 agents?

Thank you for your attention!