

Algorithms and Complexity on Temporal Graphs

George B. Mertzios

Department of Computer Science,
Durham University, UK

GROW 2017

Fields Institute
University of Toronto

October 2017

Static and Temporal Graphs

Modern networks are **highly dynamic**:

- **Social networks**: **friendships** are added/removed, individuals **leave**, new ones **enter**
- **Transportation networks**: transportation units change with time their **position** in the network
- **Physical systems**: e.g. systems of **interacting particles**

The common characteristic in all these applications:

- the **graph topology** is subject to **discrete changes** over time
- ⇒ the notion of **vertex adjacency** must be appropriately re-defined (by introducing the **time dimension** in the graph definition)

Various graph concepts (e.g. reachability, connectivity):

- crucially **depend** on the **exact temporal ordering** of the **edges**

Temporal graphs

Formally:

Definition (Temporal Graph)

A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

Temporal graphs

Formally:

Definition (Temporal Graph)

A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
 - $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.
-
- If $t \in \lambda(e)$ then edge e is **available** at time t
 - This formal definition (for **single-availabilities** per edge) embarks from:
[Kempe, Kleinberg, Kumar, *STOC*, 2000]
[Berman, *Networks*, 1996]
 - In general every edge can have **multiple availabilities**
[Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013]

Temporal graphs


Formally:

Definition (Temporal Graph)

A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

temporal graph:



temporal instances:



Temporal graphs


Formally:

Definition (Temporal Graph)


A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

temporal graph:



temporal instances:



Temporal graphs


Formally:

Definition (Temporal Graph)


A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

temporal graph:



temporal instances:



Temporal graphs


Formally:

Definition (Temporal Graph)


A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

temporal graph:



temporal instances:



Temporal graphs


Formally:

Definition (Temporal Graph)


A **temporal graph** is a pair (G, λ) where:

- $G = (V, E)$ is an **underlying (di)graph** and
- $\lambda : E \rightarrow 2^{\mathbb{N}}$ is a discrete **time-labeling** function.

temporal graph:



temporal instances:



Temporal graphs

Temporal graphs were studied under various different names:

- **time-varying** graphs
[Aaron et al., *WG*, 2014]
[Flocchini et al., *ISAAC*, 2009]
[Tang et al., *ACM Comp. Comm. Review*, 2010]
- **evolving** graphs (usually “graph-centric”)
[Avin et al., *ICALP*, 2008]
[Clementi et al., *SIAM J. Discr. Math.*, 2010]
[Ferreira, *IEEE Network*, 2004]
[Bui Xuan et al., *Int. J. Found. Comp. Sci.*, 2003]
- **dynamic** graphs
[Giakkoupis et al., *ICALP*, 2014]
[Casteigts et al., *Int. J. Par., Emergent & Distr. Syst*, 2012]
[Bhadra and Ferreira, *ADHOC-NOW*, 2003]
- **graphs over time**
[Leskovec et al., *ACM Trans. Knowl. Disc. from Data*, 2007]

Temporal graphs

Recent surveys and books:

- **Time-Varying Graphs and Dynamic Networks**
[Casteigts et al., *Int. J. Par., Emergent & Distr. Syst.*, 2012]
 - an attempt to **integrate** and **unify** existing models and concepts
- **Temporal Networks** [Holme, Saramäki, eds., *Springer*, 2013]
 - temporal network methods for **complex networks**
- **Deterministic Algorithms in Dynamic Networks**
[Casteigts, Flocchini, *Defence R&D Canada, Tech. Report I*, 2013]
[Casteigts, Flocchini, *Defence R&D Canada, Tech. Report II*, 2013]
 - survey of deterministic algorithms for **distributed computing**
 - **temporal graph classes** based on **temporal patterns** of the labels
 - satellites → periodic availabilities
 - sensor networks → connected at every instant
 - contacts in a company → bounded edge recurrence (every week)
 - community contacts → unbounded, yet recurrent interactions

Overview

- Temporal paths
- Strongly connected components
- Menger's theorem
- Temporal design problems
- Future research directions

Temporal paths

The most natural known **temporal notion** in temporal graphs:

Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, \ell_1), (e_2, \ell_2), \dots, (e_k, \ell_k))$, where:

$$\ell_1 < \ell_2 < \dots < \ell_k$$

and $\ell_i \in \lambda(e_i)$, $1 \leq i \leq k$.

Temporal paths

The most natural known **temporal notion** in temporal graphs:

Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, \ell_1), (e_2, \ell_2), \dots, (e_k, \ell_k))$, where:

$$\ell_1 < \ell_2 < \dots < \ell_k$$

and $\ell_i \in \lambda(e_i)$, $1 \leq i \leq k$.

Intuition:

- information “**flows**” along edges whose labels respect **time ordering**

Most known temporal graph parameters are “**temporal path**”-related:

- temporal versions of **distance**, **diameter**, **connectivity**, **reachability**, **exploration**, **centrality** measures, etc.

Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, \ell_1), (e_2, \ell_2), \dots, (e_k, \ell_k))$, where:

$$\ell_1 < \ell_2 < \dots < \ell_k$$

and $\ell_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:


Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, l_1), (e_2, l_2), \dots, (e_k, l_k))$, where:

$$l_1 < l_2 < \dots < l_k$$

and $l_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Temporal paths

The most natural known **temporal notion** in temporal graphs:



Definition (Temporal path; Time-respecting path; Journey)

Let (G, λ) be a **temporal graph** and $P = (e_1, e_2, \dots, e_k)$ be a walk in G . A **temporal path** is a sequence $((e_1, \ell_1), (e_2, \ell_2), \dots, (e_k, \ell_k))$, where:

$$\ell_1 < \ell_2 < \dots < \ell_k$$

and $\ell_i \in \lambda(e_i)$, $1 \leq i \leq k$.

- A **non-temporal path**:



Metrics to optimize

Question: What is the temporal analogue of an s - t shortest path?

Answer: Not uniquely defined!

- topologically shortest path: smallest number of edges
- fastest path: smallest duration
- foremost path: smallest arrival time


Metrics to optimize

Question: What is the **temporal analogue** of an **s - t shortest path**?

Answer: Not uniquely defined!

- **topologically shortest** path: smallest **number of edges**
- **fastest** path: smallest **duration**
- **foremost** path: smallest **arrival time**

Example:




Metrics to optimize

Question: What is the **temporal analogue** of an **s - t shortest path**?

Answer: Not uniquely defined!

- **topologically shortest** path: smallest **number of edges**
- **fastest** path: smallest **duration**
- **foremost** path: smallest **arrival time**

Example:



shortest: s - c - t (two edges)


Metrics to optimize

Question: What is the **temporal analogue** of an ***s-t* shortest path**?

Answer: Not uniquely defined!

- **topologically shortest** path: smallest **number of edges**
- **fastest** path: smallest **duration**
- **foremost** path: smallest **arrival time**

Example:



shortest: $s \rightarrow c \rightarrow t$ (two edges)

fastest: $s \rightarrow d \rightarrow e \rightarrow t$ (no intermediate waiting)


Metrics to optimize

Question: What is the **temporal analogue** of an ***s-t* shortest path**?

Answer: Not uniquely defined!

- **topologically shortest** path: smallest **number of edges**
- **fastest** path: smallest **duration**
- **foremost** path: smallest **arrival time**

Example:



shortest: $s-c-t$ (two edges)

fastest: $s-d-e-t$ (no intermediate waiting)

foremost: $s-a-b-t$ (arriving at time 5)

Metrics to optimize

An easy algorithm for computing all foremost paths from a given source s :
[Akrida, Gasieniec, Mertzios, Spirakis, WAOA, 2015]

- first sort the time-labels non-decreasingly
- run a BFS-like search starting from s
- at every time-step t consider only edges currently available
- if you reach a new vertex at time t , keep its predecessor

Metrics to optimize

An easy algorithm for computing **all foremost paths** from a given **source s** :
[Akrida, Gasieniec, Mertzios, Spirakis, WAOA, 2015]

Algorithm 1 Foremost Temporal Paths from Source s

```
1: Let  $S$  be the array with the sorted time-labels
2:  $R \leftarrow \{s\}$ 
3: for each  $v \in V \setminus \{s\}$  do
4:    $\text{pred}[v] \leftarrow \emptyset$ ;  $\text{arr}[v] \leftarrow \infty$    {Init.: Predecessor; Time Arrived}
5: for each time-label  $t \in S$  do
6:   for each edge  $e = (u, v)$  with  $t \in \lambda(e)$  do
7:     if  $u \in R$ ,  $v \notin R$ , and  $\text{arr}[u] < t$  then {we reached  $v$ }
8:        $\text{pred}[v] \leftarrow u$ ;  $\text{arr}[v] \leftarrow t$    {Predecessor; Time Arrived}
9:        $R \leftarrow R \cup \{v\}$ 
```

Metrics to optimize

An easy algorithm for computing **all foremost paths** from a given **source s** :

- easy adaptation of the static BFS algorithm
- running time $O(c(\lambda) \cdot \log(c(\lambda)))$
- due to the **sorting** of the labels

Algorithm 1 Foremost Temporal Paths from Source s

- 1: Let S be the array with the sorted time-labels
 - 2: $R \leftarrow \{s\}$
 - 3: **for** each $v \in V \setminus \{s\}$ **do**
 - 4: $\text{pred}[v] \leftarrow \emptyset$; $\text{arr}[v] \leftarrow \infty$ {Init.: **Predecessor**; **Time Arrived**}
 - 5: **for** each **time-label** $t \in S$ **do**
 - 6: **for** each **edge** $e = (u, v)$ with $t \in \lambda(e)$ **do**
 - 7: **if** $u \in R$, $v \notin R$, and $\text{arr}[u] < t$ **then** {we reached v }
 - 8: $\text{pred}[v] \leftarrow u$; $\text{arr}[v] \leftarrow t$ {**Predecessor**; **Time Arrived**}
 - 9: $R \leftarrow R \cup \{v\}$
-

Metrics to optimize

Polynomial algorithms exist also for computing:

- shortest and foremost paths [adaptations of Dijkstra's algorithm]
- fastest paths

[Bui-Xuan, Ferreira, Jarry, *Int. J. Found. Comp. Sci.*, 2003]

However: Not all “path-related” temporal problems tractable,
e.g. some temporal variations of:


- connectivity problems
- reachability problems

Overview

- Temporal paths
- Strongly connected components
- Menger's theorem
- Temporal design problems
- Future research directions


Temporal strongly connected components

- We write $u \rightsquigarrow v$ if there exists a **temporal path** from u to v
- The relation \rightsquigarrow is **not symmetric**: $u \rightsquigarrow v \not\iff v \rightsquigarrow u$




Temporal strongly connected components

- We write $u \rightsquigarrow v$ if there exists a **temporal path** from u to v
- The relation \rightsquigarrow is **not symmetric**: $u \rightsquigarrow v \not\Leftrightarrow v \rightsquigarrow u$




- and **not transitive**: $u \rightsquigarrow z, z \rightsquigarrow v \not\Leftrightarrow u \rightsquigarrow v$




\Rightarrow the time dimension creates its own “level of direction”

Temporal strongly connected components


static:



temporal:




temporal:




Temporal strongly connected components

static:




strongly
connected


temporal:




temporal:




Temporal strongly connected components



strongly
connected




strongly
connected
component




- $\{a, b\}$: direct temporal paths between a and b


Temporal strongly connected components



strongly
connected






strongly
connected
component



strongly
connected
component




- $\{a, b\}$: direct temporal paths between a and b
- $\{b, c\}$: the only temporal path from c to b passes through $a \notin \{b, c\}$

Temporal strongly connected components



- $\{a, b\}$: direct temporal paths between a and b
- $\{b, c\}$: the only temporal path from c to b passes through $a \notin \{b, c\}$
- $\{a, b, c\}$: no temporal path from a to c

Temporal strongly connected components






Definition (Bharda, Ferreira, 2003)

An **open strongly connected component (o-SCC)** in a temporal graph is a set S of vertices such that $u \rightsquigarrow v$ for every $u, v \in S$.

Examples of an o-SCC: $\{a, b\}$, $\{b, c\}$

Temporal strongly connected components



Definition (Bharda, Ferreira, 2003)

An **open strongly connected component (o-SCC)** in a temporal graph is a set S of vertices such that $u \rightsquigarrow v$ for every $u, v \in S$.

Examples of an o-SCC: $\{a, b\}$, $\{b, c\}$

Definition (Bharda, Ferreira, 2003)


A **strongly connected component (SCC)** in a temporal graph is a set S of vertices such that, for every $u, v \in S$, there is a **temporal path** from u to v that uses **only** vertices from S .

Example of a SCC: $\{a, b\}$

Temporal strongly connected components

A further difference to the static case:


- two **different SCCs** can have **common vertices**



Temporal strongly connected components

A further difference to the static case:


- two **different SCCs** can have **common vertices**
 - $\{a, c, d\}$ is a SCC



Temporal strongly connected components

A further difference to the static case:


- two **different SCCs** can have **common vertices**
 - $\{a, c, d\}$ is a SCC
 - $\{b, c, d\}$ is another SCC



Temporal strongly connected components

A further difference to the static case:

- two **different SCCs** can have **common vertices**
 - $\{a, c, d\}$ is a SCC
 - $\{b, c, d\}$ is another SCC
 - $\{a, b, c, d\}$ is **not** a SCC (no temporal path $b \rightsquigarrow a$)



Temporal strongly connected components

Theorem (Bharda, Ferreira, 2003)

*Given a vertex subset S of a temporal graph (G, λ) , we can **verify** in **polynomial time** whether S is a **SCC** (resp. an **o-SCC**).*

- Proof: similarly to static graphs

Temporal strongly connected components

However:

Theorem

Given a temporal graph (G, λ) , it is *NP-hard* to compute the *maximum* size of a *SCC*, even if all edges have *one* and the *same label*.

Temporal strongly connected components


However:

Theorem

Given a temporal graph (G, λ) , it is **NP-hard** to compute the **maximum size of a SCC**, even if all edges have **one** and the **same label**.

Proof (sketch):

- Reduction from CLIQUE.
- Given a **static undirected** graph G construct a **temporal** graph (\mathcal{G}, λ) :
 - 1 for each v_i of G create vertex u_i in \mathcal{G} ,
 - 2 for each **edge** (v_i, v_j) of G add these **two arcs** to \mathcal{G} :



- G has a **clique** of size $k \Leftrightarrow (\mathcal{G}, \lambda)$ has an **SCC** of size k .



Overview

- Temporal paths
- Strongly connected components
- Menger's theorem
- Temporal design problems
- Future research directions

Menger's Theorem

Two fundamental **duality** results in (static) graph theory:

Theorem (Menger, 1927)


*The **maximum** number of **vertex-disjoint** (**edge-disjoint**) **s-t paths** is equal to the **minimum** number of **vertices** (**edges**) needed to **separate s** from **t**.*

Menger's Theorem

Two fundamental **duality** results in (static) graph theory:

Theorem (Menger, 1927)

The **maximum** number of **vertex-disjoint** (**edge-disjoint**) **s - t paths** is equal to the **minimum** number of **vertices** (**edges**) needed to **separate** s from t .




Menger's Theorem

Two fundamental **duality** results in (static) graph theory:

Theorem (Menger, 1927)

The **maximum** number of **vertex-disjoint** (**edge-disjoint**) **s - t paths** is equal to the **minimum** number of **vertices** (**edges**) needed to **separate** s from t .




Menger's Theorem

Two fundamental **duality** results in (static) graph theory:

Theorem (Menger, 1927)

The **maximum** number of **vertex-disjoint** (**edge-disjoint**) **s - t paths** is equal to the **minimum** number of **vertices** (**edges**) needed to **separate** **s** from **t** .




Menger's Theorem

A temporal analogue of the “edge-version”:

Theorem (Berman, 1996)

In *single-labeled* temporal graphs, the *maximum* number of *edge-disjoint temporal s - t paths* is equal to the *minimum* number of *edges* needed to *temporally separate s from t* .




Menger's Theorem

A temporal analogue of the “edge-version”:

Theorem (Berman, 1996)

In *single-labeled* temporal graphs, the *maximum* number of *edge-disjoint temporal s - t paths* is equal to the *minimum* number of *edges* needed to *temporally separate s from t* .



Menger's Theorem

However the “intuitive” temporal **vertex**-version **fails**:

Lemma (Berman, 1996; Kempe, Kleinberg, Kumar, 2000)

There exists a single-labeled temporal graph where:

***maximum** number of **vertex**-disjoint **temporal** s - t paths $<$*

***minimum** number of **vertices** needed to **temporally separate** s from t .*


Menger's Theorem

However the “intuitive” temporal **vertex**-version **fails**:

Lemma (Berman, 1996; Kempe, Kleinberg, Kumar, 2000)

There exists a single-labeled temporal graph where:
maximum number of vertex-disjoint temporal s - t paths $<$
minimum number of vertices needed to temporally separate s from t .

- no two vertex-disjoint temporal s - t paths




Menger's Theorem

However the “intuitive” temporal **vertex**-version **fails**:

Lemma (Berman, 1996; Kempe, Kleinberg, Kumar, 2000)

There exists a single-labeled temporal graph where:
maximum number of vertex-disjoint temporal s - t paths $<$
minimum number of vertices needed to temporally separate s from t .

- no two vertex-disjoint temporal s - t paths
- the removal of any one vertex leaves s and t temporally connected




Menger's Theorem

However the “intuitive” temporal **vertex**-version **fails**:

Lemma (Berman, 1996; Kempe, Kleinberg, Kumar, 2000)

There exists a single-labeled temporal graph where:
maximum number of vertex-disjoint temporal s - t paths $<$
minimum number of vertices needed to temporally separate s from t .

- no two vertex-disjoint temporal s - t paths
- the removal of any one vertex leaves s and t temporally connected




Menger's Theorem

However the “intuitive” temporal **vertex**-version **fails**:

Lemma (Berman, 1996; Kempe, Kleinberg, Kumar, 2000)

There exists a single-labeled temporal graph where:
maximum number of vertex-disjoint temporal s - t paths $<$
minimum number of vertices needed to temporally separate s from t .

- no two vertex-disjoint temporal s - t paths
- the removal of any one vertex leaves s and t temporally connected



Menger's Theorem

An appropriate temporal vertex-version of Menger's theorem

We say that:

- two temporal paths are **out-disjoint** if they never **leave** from the **same node** at the **same time**
- we remove **departure time t** from **vertex u** if:
 - we remove **label t** from for all edges **(u, w)**

Menger's Theorem

An appropriate temporal vertex-version of Menger's theorem

We say that:

- two temporal paths are **out-disjoint** if they never **leave** from the **same node** at the **same time**
- we remove **departure time t** from **vertex u** if:
 - we remove **label t** from for all edges **(u, w)**

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, ICALP 2013)


*In **multi-labeled** temporal graphs, the **maximum** number of **out-disjoint** s - t temporal paths equals the **minimum** number of **vertex departure times** needed to temporally separate s from t .*

- **vertex** disjointness \longrightarrow **vertex departure** time disjointness
- **vertex** removal \longrightarrow **vertex departure** time removal

Menger's Theorem

An appropriate temporal vertex-version of Menger's theorem

Three out-disjoint temporal paths from s to t :



Overview

- Temporal paths
- Strongly connected components
- Menger's theorem
- Temporal design problems
- Future research directions

Temporal design problems

In many scheduling problems:

- the provided **graph topology** G represents a given **static specification**
 - e.g. available **bus routes** in the city center
- the aim is to organize a **temporal schedule** on this specification, e.g.
 - **when** the buses should be in **which** stop
 - such that **every pair** of stops is **connected** via a route
- while minimizing some **cost function**
 - e.g. with as **few buses** as possible

Temporal design problems

We mainly study the following **cost functions** of a **time-label λ** :

- ① **temporality τ** : the maximum number of labels per edge
 - a **distributed / decentralized** measure of cost in the temporal network
- ② **temporal cost κ** : the total number of labels on all edges
 - a **centralized** measure of cost

Temporal design problems

We mainly study the following **cost functions** of a **time-label** λ :

- 1 **temporality** τ : the maximum number of labels per edge
 - a **distributed / decentralized** measure of cost in the temporal network
- 2 **temporal cost** κ : the total number of labels on all edges
 - a **centralized** measure of cost

and two fundamental **connectivity properties**:


- 1 **preserve** in (G, λ) **all reachabilities** in G
 - if v is **reachable** from u in $G \Rightarrow u \rightsquigarrow v$ in (G, λ)
- 2 **preserve** in (G, λ) **all paths** in G
 - G has a **path** $P \Rightarrow (G, \lambda)$ has a **temporal path** on the **same edges** as P

Temporality of the ring C_n

- The labeling assigning to each edge (u_i, u_{i+1}) the labels $\{i, n + i\}$ preserves all paths, i.e. $\tau(C_n, \text{all paths}) \leq 2$

$$\Rightarrow \tau(C_n, \text{all paths}) = 2$$

- The maximum label is $2n$
(can be “tuned” to $2n - 2$)



Temporality of the ring C_n

- The labeling assigning to each edge (u_i, u_{i+1}) the labels $\{i, n + i\}$ preserves all paths, i.e. $\tau(C_n, \text{all paths}) \leq 2$

$$\Rightarrow \tau(C_n, \text{all paths}) = 2$$


- The maximum label is $2n$
(can be “tuned” to $2n - 2$)

What if we restrict the age to $\alpha(\lambda) = n - 1$?

- Assume that some edge e misses label $i \in \{1, 2, \dots, n - 1\}$
- Then there exists a temporal path on C_n that needs label i on edge e to finish by time $n - 1$

\Rightarrow the optimal labeling assigns $\{1, 2, \dots, n - 1\}$ to all edges of C_n


$$\Rightarrow \tau(C_n, \text{all paths}, n - 1) = n - 1$$



Temporality of a DAG

Lemma (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

If G is a **DAG** then $\tau(G, \text{all paths}) = 1$.




Temporality of a DAG

Lemma (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

If G is a **DAG** then $\tau(G, \text{all paths}) = 1$.

Proof.

- Take a **topological sort** u_1, u_2, \dots, u_n of G
- Give **label** i to every **edge** (u_i, u_j) , where $i < j$.



A generic method for lower bounds of temporality

Intuition gained:

- cycles can increase the temporality of a (di)graph

A generic method for lower bounds of temporality


Intuition gained:

- cycles can increase the temporality of a (di)graph

Based on this intuition:

Definition (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

The set $S = \{e_1, e_2, \dots, e_k\} \subseteq E(G)$ is an **edge-kernel** of G if for every permutation of S there is a (static) path of G that visits all edges of S according to this permutation.



A generic method for lower bounds of temporality

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

If a (di)graph G has an *edge-kernel* of size k then $\tau(G, \text{all paths}) \geq k$.

A generic method for lower bounds of temporality

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

If a (di)graph G has an *edge-kernel* of size k then $\tau(G, \text{all paths}) \geq k$.

Lemma (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

The *complete (di)graph* on n vertices has an *edge-kernel* of size $\lfloor n/2 \rfloor$.

A generic method for lower bounds of temporality

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

If a (di)graph G has an *edge-kernel* of size k then $\tau(G, \text{all paths}) \geq k$.

Lemma (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

The *complete (di)graph* on n vertices has an *edge-kernel* of size $\lfloor n/2 \rfloor$.

Lemma (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

There exist *(undirected) planar graphs* with an *edge-kernel* of size $\Omega(n^{\frac{1}{3}})$.

Temporality: preserving all reachabilities

If we want to preserve only the reachabilities:

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

Let G be an *undirected* (or *strongly connected directed*) graph.
Then $\tau(G, \text{reach}) \leq 2$.

Proof idea:

- pick an arbitrary vertex v (as the “root”)
- build a temporal *in-tree* to vertex v
- from v build a temporal *out-tree* to vertex v

Temporality: preserving all reachabilities

If we want to preserve only the reachabilities:

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

Let G be an *undirected* (or *strongly connected directed*) graph.
Then $\tau(G, \text{reach}) \leq 2$.

Proof idea:

- pick an arbitrary vertex v (as the “root”)
- build a temporal *in-tree* to vertex v
- from v build a temporal *out-tree* to vertex v

Similarly to our analysis for DAGs:

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

Let G be a *directed* graph. Then $\tau(G, \text{reach}) = \max_{C \in \mathcal{C}(G)} \tau(C, \text{reach})$,
where $\mathcal{C}(G)$ is the set of *strongly connected components* of G .

Temporal cost for preserving all reachabilities

A very different cost function: total **number κ of labels**

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, *ICALP*, 2013)

Let $d(G)$ denote the (static) **diameter** of the directed graph G . The problem of computing $\kappa(G, \text{reach}, d(G))$ is **APX-hard**.

Temporal cost for preserving all reachabilities

A very different cost function: total **number κ of labels**

Theorem (Mertzios, Michail, Chatzigiannakis, Spirakis, ICALP, 2013)

Let $d(G)$ denote the (static) **diameter** of the directed graph G . The problem of computing $\kappa(G, \text{reach}, d(G))$ is **APX-hard**.

The “inverse” design problem:

- given a **temporal graph** (G, λ) that maintains **all reachabilities** of G
- **remove** the maximum number of labels by **maintaining reachabilities**
- **removal cost** $r(G, \lambda)$

Theorem (Akrida, Gasieniec, Mertzios, Spirakis, WAOA, 2015)

The problem of computing $r(G, \lambda)$ is **APX-hard** on undirected graphs G .

Overview

- Temporal paths
- Strongly connected components
- Menger's theorem
- Temporal design problems
- Future research directions

- Find **constant-factor** approximations for the various temporal graph design problems
- Other natural **connectivity properties** subject to which optimization is to be performed
- Efficient **deterministic/randomized/approximation** algorithms on special **temporal graph classes**, i.e. by restricting:
 - the underlying **topology G** and/or
 - the **temporal pattern** with which the time-labels appear (a new dimension with no previous static analogue!)

- Temporal graphs defined by the mobility patterns of mobile wireless entities modeled by a **sequence of unit disk graphs**
 - **Well-motivated** as a natural source of temporal graphs
 - May allow for **better approximations**
- Other natural **non-path temporal problems** (apart from matchings)
 - a recently defined notion of a “ Δ -temporal clique” in social networks:
“a set of nodes and a time interval such that all pairs interact at least every Δ during this interval”
[Viard, Latapy, Magnien, *ASONAM*, 2015]
[Himmel, Molter, Niedermeier, Sorge, *Social Network Analysis and Mining*, 2017]
- Our results so far are a first step towards answering this fundamental question:

To what extent can algorithmic and structural results of graph theory be carried over to temporal graphs?

Research Directions

New EPSRC project
(Starting September 2017):

Title: **Algorithmic Aspects of Temporal Graphs**

P.I. in **Durham**: George Mertzios

- PostDoc: Viktor Zamaraev

P.I. in **Liverpool**: Paul Spirakis

- PostDoc: Eleni Akrida

Value total amount: £ 800,000

Thank you for your attention!